

Data vault method coverage

This document provides examples for all expected or seen model and loading scenarios. Please note, that some of the syntax is only announced but not yet implemented in the compiler reference implementation.

Also please keep in mind, that json does not use linefeed or indentation as syntax elements. **The examples are formatted for better human readability only.**

Credits and license

(C) Matthias Wegner, cimt ag

Creative Commons License [CC BY-ND 4.0](#)

This concept is published under the Creative Commons License CC BY-ND 4.0.

It allows reusers to copy and distribute the material in any medium or format in **unadapted form only**, and only so long as **attribution is given to the creator**. The license allows for commercial use. So incorporating the concept into a commercial product is allowed.

Data Vault stereo type coverage and examples

The following examples provide field and model DVPD declarations for all official data vault stereotypes and their variations. It includes all examples to the stereotypes, explained in the book "Building A Scalable Data Warehouse With Data Vault 2.0" by Dan Linstedt and Michael Olschimke from 2016. These examples are marked with "{DV-<chapter>.<subchapter>}" in the heading.

To be easy understandable, the examples use simplified table and column names that don't follow all best practices. For the same reason, not all examples are matched to the data vault book.

Some properties of the DVPD can be declared on the level of the **model profile**. All examples refer to the following, most common, profile settings:

```
{  "compare_criteria": "key+actual",
  "uses_diff_hash_default": true,
  "is_enddated_default": true,
  "has_deletion_flag_default": true
}
```

Hub tables

Hub with a single column business key {DV-4.3}

```
"fields": [
  {  "field_name": "AIRLINE_ID"
```

```

        , "field_type": "Varchar(20)"
        , "targets": [{"table_name": "airline_hub"}]
    },
    ...
],
"tables": [
    {
        "table_name": "airline_hub"
        , "table_stereotype": "hub"
        , "hub_key_column_name": "HK_AIRLINE"
    }
    ...
]

```

Hub with a composite business key {DV-4.3.1.1}

This example assumes, that customer id's are not unique over the different web shops, the data is collected from. Therefore the web shop id must also be used for identification.

```

"fields": [
    {
        "field_name": "WEBSHOP_ID",
        "field_type": "integer",
        "targets": [{"table_name": "customer_hub"}]
    },
    {
        "field_name": "CUSTOMER_ID",
        "field_type": "Varchar(20)",
        "targets": [{"table_name": "customer_hub"}]
    },
    ...
],
"tables": [
    {
        "table_name": "customer_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_CUSTOMER"
    },
    ...
]

```

Hub with a last seen date {DV-4.3.2.5}

```

"fields": [
    {
        "field_name": "TAILNUM",
        "field_type": "Varchar(20)",
        "targets": [{"table_name": "airline_hub"}]
    },
    {
        "field_name": "LASTSEENDATE",
        "field_type": "TIMESTAMP",
        "field_value": "${LOAD_TIMESTAMP}",
        "targets": [
            {"table_name": "airline_hub"}
        ]
    }
]

```

```

        , "exclude_from_key_hash":true
        , "update_on_every_load":true}
    ]
  },
  ...
],
"tables": [
  {
    "table_name": "airplane_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_AIRPLANE"
  }
  ...
]

```

Multiple hubs with partly the same business key

Here for both hub's the business key must contain the company id to be unique. Also the incoming field ARCHITECT contains an employee identification must be mapped to the business key column EMPLOYEE_ID

```

"fields": [
  {
    "field_name": "COMPANY_ID",
    "field_type": "Varchar(20)",
    "targets": [ {"table_name": "employee_hub"},
                  {"table_name": "building_hub"} ]
  },
  {
    "field_name": "ARCHITECT",
    "field_type": "integer",
    "targets": [ {"table_name": "employee_hub",
                    "column_name": "EMPLOYEE_ID"
                  } ]
  },
  {
    "field_name": "BUILDING_SIGNATURE",
    "field_type": "Varchar(100)",
    "targets": [ {"table_name": "building_hub"} ]
  }
  ...
],
"tables": [
  {
    "table_name": "employee_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_EMPLOYEE"
  }
  {
    "table_name": "building_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_BUILDING"
  }
  ...
]

```

Link tables

Link connecting two hubs {DV-4.3}

An order and its related customer

```
"fields": [
  {
    "field_name": "ORDER_ID",
    "field_type": "integer",
    "targets": [ {"table_name": "order_hub"} ]
  },
  {
    "field_name": "CUSTOMER_ID",
    "field_type": "Varchar(20)",
    "targets": [ {"table_name": "customer_hub"} ]
  }
],
"tables": [
  {
    "table_name": "order_customer_link",
    "table_stereotype": "lnk",
    "link_parent_tables": ["order_hub", "customer_hub"]
  },
  {
    "table_name": "order_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_ORDER"
  },
  {
    "table_name": "customer_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_CUSTOMER"
  }
]
```

Link connecting three hubs {DV-4.4.2}

Example of a car rental relation. A rental references a customer, a car and a rental station.

```
"fields": [
  {
    "field_name": "VEHICLE_ID",
    "field_type": "integer",
    "targets": [ {"table_name": "car_hub"} ]
  },
  {
    "field_name": "CUSTOMER_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "customer_hub"}]
  },
  {
    "field_name": "STATION_NAME",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "station_hub"}]
  }
],
"tables": [
```

```

    {
      "table_name": "car_customer_station_link",
      "table_stereotype": "lnk",
      "link_parent_tables": [ "car_hub",
                              "customer_hub",
                              "station_hub" ]
    },
    {
      "table_name": "car_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_CAR"
    },
    {
      "table_name": "customer_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_CUSTOMER"
    },
    {
      "table_name": "station_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_STATION"
    },
  ]

```

Link with two hubs, but one with two relations {DV-4.4.4}

The link describes a flight, referencing the flight carrier and two airports (1 source and 1 destination).

```

"fields": [
  {
    "field_name": "CARRIER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "carrier_hub"}]
  },
  {
    "field_name": "SOURCEAIRPORT_ID",
    "field_type": "Varchar(20)",
    "targets": [
      {
        "table_name": "airport_hub",
        "column_name": "AIRPORT_ID",
        "relation_name": "SOURCE"
      }
    ]
  },
  {
    "field_name": "DESTINATIONAIRPORT_ID",
    "field_type": "Varchar(20)",
    "targets": [
      {
        "table_name": "airport_hub",
        "column_name": "AIRPORT_ID",
        "relation_name": "DEST"
      }
    ]
  },
  ...
],
"tables": [
  {
    "table_name": "car_customer_station_link",
    "table_stereotype": "lnk",
    "link_parent_tables": [
      { "table_name": "carrier_hub",
        "relation_name": "SOURCE" },
      { "table_name": "airport_hub",
        "relation_name": "SOURCE" }
    ]
  }
]

```

```

        }
        {"table_name": "airport_hub",
         "relation_name": "DEST"
        } ]
    },
    {
        "table_name": "carrier_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_CARRIER"
    },
    {
        "table_name": "airport_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_AIRPORT"
    },
    ...
]

```

Link with dependent child key columns {DV-4.4.5.2}

Dependent child keys are declared by mapping fields to link tables. In this case the selling month and year are placed as dependent child keys of the webshop sale report link.

```

"fields": [
    {
        "field_name": "WEBSHOP_ID",
        "field_type": "integer",
        "targets": [{"table_name": "webshop_hub"}]
    },
    {
        "field_name": "PRODUCT_ID",
        "field_type": "integer",
        "targets": [{"table_name": "product_hub"}]
    },
    {
        "field_name": "SELLING_MONTH",
        "field_type": "integer",
        "targets": [{"table_name": "webshop_sale_report_link"}]
    },
    {
        "field_name": "SELLING_YEAR",
        "field_type": "integer",
        "targets": [{"table_name": "webshop_sale_report_link"}]
    },
    ...
],
"tables": [
    {
        "table_name": "webshop_sale_report_link",
        "table_stereotype": "lnk",
        "link_parent_tables": ["webshop_hub", "product_hub"]
    },
    {
        "table_name": "webshop_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_WEBSHOP"
    },
    {
        "table_name": "product_hub",
        "table_stereotype": "hub",

```

```

        "hub_key_column_name": "HK_PRODUCT"
    }
]

```

Link with last seen date {DV-4.4.5}

An order and its related customer. The last seen date is the timestamp of the load process. Its not part of the source, but added by the loading process. It is stored in the link, not relevant for the link hash key and must be updated every time. This behaviour is declared in the field and its mapping.

```

"fields": [
  {
    "field_name": "ORDER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "order_hub"}]
  },
  {
    "field_name": "CUSTOMER_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "customer_hub"}]
  },
  {
    "field_name": "LASTSEENDATE",
    "field_type": "TIMESTAMP",
    "field_value": "${LOAD_TIMESTAMP}",
    "targets": [{
      "table_name": "order_customer_link",
      "exclude_from_key_hash": true,
      "update_on_every_load": true
    }]
  },
],
"tables": [
  {
    "table_name": "order_customer_link",
    "table_stereotype": "lnk",
    "link_parent_tables": ["order_hub", "customer_hub"]
  },
  {
    "table_name": "order_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_ORDER"
  },
  {
    "table_name": "customer_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_CUSTOMER"
  }
]

```

Same As Link {DV-5.1.1}{DV-5.2.2}

The source for a "same as" Link contains the business keys of the main object and the business key of its duplicate. The following example loads a "Same as Link" that contains one hub key HK_PRODUCT for the master product and HK_PRODUCT_DUPLICATE for the duplicate.

```

"fields": [
  {
    "field_name": "PRODUCT_ID",
    "field_type": "integer",
    "targets": [{"table_name": "product_hub"}]},
  {
    "field_name": "SAME_PRODUCT_ID",
    "field_type": "integer",
    "targets": [{
      "table_name": "product_hub",
      "column_name": "PRODUCT_ID",
      "relation_name": "DUPLICATE"
    }]
  },
],
"tables": [
  {
    "table_name": "product_duplicate_saslink",
    "table_stereotype": "lnk",
    "link_parent_tables": [
      {
        "table_name": "product_hub",
        {
          "table_name": "product_hub",
          "relation_name": "DUPLICATE"
        }
      }
    ],
  },
  {
    "table_name": "product_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_PRODUCT"
  }
]

```

To also load the SAME_PRODUCT_ID id to the product hub, its field mapping needs to declare an explicit relation name, to distinguish it from PRODUCT_ID. It also needs to declare the column name to be PRODUCT_ID. The same relation name must be declared in the mapping of link parent tables. This determines, in which hub key column in the link with hash value of must be placed. The link will generate its own hub key column name for the second column, from the original name and the name of the relation (if not told otherwise).

Link on Link {DV-5.2.1}

A link on a link is **not supported** by the core syntax of DVPD, since it is highly discouraged by Dan Linstedt. The Data Vault methodology already describes ways to circumvent the need for it.

hierarchical link {DV-5.2.3}

This is an example of a part containment hierarchy. The declaration is another variety of multiple parent relations to the same hub.

```

"fields": [
  {
    "field_name": "PART_ID",
    "field_type": "integer",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "CONTAINING_PART_ID",

```



```

        "field_type": "integer",
        "targets": [
            {
                "table_name": "product_hub",
                "column_name": "PART_ID",
                "relation_name": "CONTAINED_BY"
            }
        ],
    },
],
"tables": [
    {
        "table_name": "product_containment_hlink",
        "table_stereotype": "lnk",
        "link_parent_tables": [
            { "table_name": "product_hub" },
            { "table_name": "product_hub", "relation_name": "CONTAINED_BY" }
        ],
        {
            "table_name": "product_hub",
            "table_stereotype": "hub",
            "hub_key_column_name": "HK_PRODUCT"
        }
    ]
]

```

non historized link {DV-5.2.4}

```

"fields": [
    {
        "field_name": "ACCOUNT_NO",
        "field_type": "integer",
        "targets": [ { "table_name": "account_hub" } ]
    },
    {
        "field_name": "ACCOUNTANT_ID",
        "field_type": "varchar(20)",
        "targets": [ { "table_name": "accountant_hub" } ]
    },
    {
        "field_name": "BOOKING_ID",
        "field_type": "varchar(22)",
        "targets": [ { "table_name": "account_booking_tlink" } ]
    },
    {
        "field_name": "BOOKING_TIME",
        "field_type": "varchar(22)",
        "targets": [ { "table_name": "account_booking_tlinksat" } ]
    },
    {
        "field_name": "AMOUNT",
        "field_type": "decimal(12,2)",
        "targets": [ { "table_name": "account_booking_tlinksat" } ]
    },
],
"tables": [
    {
        "table_name": "account_booking_tlink",
        "table_stereotype": "lnk",
        "link_parent_tables": [ "account_hub", "accountant_hub" ]
    }
]

```

```

    },
    {
      "table_name": "account_booking_tlinksat",
      "table_stereotype": "sat",
      "satellite_parent_table": "account_booking_tlink",
      "is_enddated": false,
      "has_deletion_flag": false,
      "compare_criteria": "key"
    }
  },
  {
    "table_name": "account_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_ACCOUNT"
  },
  {
    "table_name": "accountant_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_ACCOUNTANT"
  }
]

```

nondescriptive link (single) {DV-5.2.5}

Like the simple link example, non descriptive links are expressed by leaving out satellites. To keep only the current relations and remove deprecated relations a deletion rule for the link is declared. The source delivers complete customer relations for all orders in the data increment, so the order is partition criteria.

This example focuses on only one link from the books Figure 5.17. Since the DVPD core syntax only supports 1 stage table, loading both links in Figure 5.17 of the book, needs two pipelines.

```

"fields": [
  {
    "field_name": "OFFERING_ID",
    "field_type": "integer",
    "targets": [{"table_name": "offering_hub"}]
  },
  {
    "field_name": "CUSTOMER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "customer_hub"}]
  }
],
"tables": [
  {
    "table_name": "customer_offering_interest_ndlink",
    "table_stereotype": "lnk",
    "link_parent_tables": ["offering_hub", "customer_hub"]
  },
  {
    "table_name": "offering_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_OFFERING"
  },
  {
    "table_name": "customer_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_CUSTOMER"
  }
]

```

```

    ]
    "deletion_detection": {
        "procedure": "stage_comparison",
        "partitioning_fields": ["OFFERING_ID"],
        "deletion_rules": [
            {
                "tables_to_cleanup": ["customer_offering_interest_ndlink"],
                "join_path": ["offering_hub"]
            }
        ]
    }
}

```

By restricting the deletion detection to the current OFFERING_ID in the stage, incremental data delivery will not remove relations of offerings, that are not in the current data set.

nondescriptive link (multi) {DV-5.2.5}

needs review <<

```

"fields": [
    {
        "field_name": "OFFERING_ID",
        "field_type": "integer",
        "targets": [{"table_name": "offering_hub"}]
    },
    {
        "field_name": "CUSTOMER_ID",
        "field_type": "integer",
        "targets": [{"table_name": "customer_hub"}]
    },
],
"tables": [
    {
        "table_name": "customer_offering_interest_ndlink",
        "table_stereotype": "lnk",
        "link_parent_tables": ["offering_hub", "customer_hub"]
    },
    {
        "table_name": "offering_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_OFFERING"
    },
    {
        "table_name": "customer_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_CUSTOMER"
    }
]
"deletion_detection": {
    "procedure": "stage_comparison",
    "partitioning_fields": ["OFFERING_ID"],
    "deletion_rules": [
        {
            "tables_to_cleanup": ["customer_offering_interest_ndlink"]
        }
    ]
}

```

exploration link {DV-5.2.7}

An exploration link is declared like any other link by declaring the hubs, that are connected by the link and the link itself. The main difference to normal links comes from the sourcing of the business keys, that will be selected from the raw vault. (a directive to take hub key values from the source dataset instead of recalculating it, will be added in later versions)

Satellite tables

Normal Satellite on a hub {DV-4.5}

```
"fields": [
  {
    "field_name": "AIRPORTID",
    "field_type": "integer",
    "targets": [{"table_name": "airport_hub"}]
  },
  {
    "field_name": "RUNWAYLENGTH",
    "field_type": "DECIMAL(5,1)",
    "targets": [{"table_name": "airport_sat"}]
  },
  {
    "field_name": "RUNWAYELEVATION",
    "field_type": "DECIMAL(10,0)",
    "targets": [{"table_name": "airport_sat"}]
  }
],
"tables": [
  {
    "table_name": "airport_sat",
    "table_stereotype": "sat",
    "satellite_parent_table": "airport_hub",
    "diff_hash_column_name": "DIFF_AIRPORT_SAT"
  },
  {
    "table_name": "airport_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_AIRPORT"
  }
]
```

Multiple satellites on a hub (Splitting by rate of change) {DV-4.5.2.2}

Store the fast changing attributes of a product (price, priority) separate from the slow/never changing attribute (name, class). Just as reminder: A pipeline only transforms one source object. So splitting by source system is achieved by using different pipelines serving different satellites.

```
"fields": [
  {
    "field_name": "GTIN",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "PRICE",
    "field_type": "DECIMAL(10,2)",
    "targets": [{"table_name": "product_sat"}]
  }
]
```

```

    {
      "field_name": "NAME",
      "field_type": "Varchar(20)",
      "targets": [{"table_name": "product_slow_sat"}]
    },
    {
      "field_name": "CLASS",
      "field_type": "Varchar(20)",
      "targets": [{"table_name": "product_slow_sat"}]
    },
    {
      "field_name": "PRICE",
      "field_type": "Varchar(20)",
      "targets": [{"table_name": "product_fast_sat"}]
    },
    {
      "field_name": "PRIORITY",
      "field_type": "Varchar(20)",
      "targets": [{"table_name": "product_fast_sat"}]}
  ],
  "tables": [
    {
      "table_name": "product_slow_sat",
      "table_stereotype": "sat",
      "satellite_parent_table": "customer_hub",
      "diff_hash_column_name": "DIFF_PRODUCT_SLOW_SAT"
    },
    {
      "table_name": "product_fast_sat",
      "table_stereotype": "sat",
      "satellite_parent_table": "customer_hub", "diff_hash_column_name":
"DIFF_PRODUCT_FAST_SAT"
    },
    {
      "table_name": "product_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_PRODUCT"
    }
  ]

```

Satellite with extract date {DV-4.5.3.5}

In case the extract date differs significantly from the loading date, it must be somehow part of the incoming data. To prevent historization of every new extraction, when the content has not changed, the extract date is excluded from the comparison during the satellite loading.

```

"fields": [
  {
    "field_name": "AIRPORTID",
    "field_type": "integer",
    "targets": [{"table_name": "airport_hub"}]
  },
  {
    "field_name": "RUNWAYLENGTH",
    "field_type": "DECIMAL(5,1)",
    "targets": [{"table_name": "airport_sat"}]
  },
  {
    "field_name": "RUNWAYELEVATION",
    "field_type": "DECIMAL(10,0)",
    "targets": [{"table_name": "airport_sat"}]
  }
]

```

```

    }
    {
      "field_name": "EXTRACT_TIMESTAMP",
      "field_type": "TIMESTAMP",
      "targets": [{
        "table_name": "airport_sat",
        "exclude_from_change_detection": true
      }]
    }
  ],
  "tables": [
    {
      "table_name": "airport_sat",
      "table_stereotype": "sat",
      "satellite_parent_table": "airport_hub",
      "diff_hash_column_name": "DIFF_AIRPORT_SAT"
    },
    {
      "table_name": "airport_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_AIRPORT"
    }
  ]

```

Satellite on a link, with a driving key declaration {DV-4.5.5}

This data source is a table with the order, referencing the product of the order. Should product of the order be modified the former product of the order must be "unlinked". Data Vault indicates this by declaring the driving keys, that must be used by the loading process for ending former relations. Driving keys are the hub key columns of the parent link of the satellite.

```

"fields": [
  {
    "field_name": "ORDER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "order_hub"}]
  },
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "PRICE",
    "field_type": "DECIMAL(16,2)",
    "targets": [{"table_name": "order_product_sale_sat"}]
  },
  {
    "field_name": "QUANTITY",
    "field_type": "DECIMAL(8,0)",
    "targets": [{"table_name": "order_product_sale_sat"}]
  },
],
"tables": [
  {
    "table_name": "order_product_sale_sat",
    "table_stereotype": "sat",
    "satellite_parent_table": "order_product_link",
    "driving_keys": ["HK_ORDER"],
    "diff_hash_column_name": "DIFF_ORDER_PRODUCT_SALES_SAT"
  }
]

```

```

    },
    {
      "table_name": "order_product_link",
      "table_stereotype": "lnk",
      "link_parent_tables": ["order_hub", "product_hub"]
    },
    {
      "table_name": "order_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_ORDER"
    },
    {
      "table_name": "product_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_PRODUCT"
    }
  ]

```

Multi-Active Satellite {DV-5.3.2}

The example shows a possible solution, how to store the matching scores of products to various categories in a multi active satellite. The only declaration difference to a normal satellite is the "is_multiactive" attribute that should trigger the specific load processing of multi active satellite.

```

"fields": [
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "CATEGORY",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_msat"}]
  },
  {
    "field_name": "MATCH_SCORE",
    "field_type": "integer",
    "targets": [{"table_name": "product_msat"}]
  }
],
"tables": [
  {
    "table_name": "product_msat",
    "table_stereotype": "sat",
    "is_multiactive": "true",
    "satellite_parent_table": "product_hub",
    "diff_hash_column_name": "DIFF_PRODUCT_MSAT"
  },
  {
    "table_name": "product_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_PRODUCT"
  }
]

```

Status Tracking Satellite with sequence information {Dv-5.3.3}

In this example, the CDC information contains a change timestamp from the source system. This is used to prevent reloading CDC information, already received. We skip the creation of a diff hash, since the data is very small.

```
"fields": [
  {
    "field_name": "EMPLOYEE_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "employee_hub"}]
  },
  {
    "field_name": "CHANGE_TIMESTAMP",
    "field_type": "timestamp",
    "targets": [{"table_name": "employee_tracksat"}]
  },
  {
    "field_name": "OPERATION",
    "field_type": "Varchar(3)",
    "targets": [{"table_name": "employee_tracksat"}]
  }
],
"tables": [
  {
    "table_name": "employee_tracksat",
    "table_stereotype": "sat",
    "satellite_parent_table": "employee_hub", "uses_diff_hash": "false"
  },
  {
    "table_name": "employee_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_EMPLOYEE"
  }
]
```

Tracking satellite without sequence information

In this example, the CDC information contains no data about the sequence of events. Incoming events must always be stored. The load timestamp in the satellite will be the only indicator about the sequence of events. Duplication of data by reloading the same events must be prevented by the loading process.

```
"fields": [
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "OPERATION",
    "field_type": "Varchar(3)",
    "targets": [{"table_name": "product_tracksat"}]
  }
],
"tables": [
  {
    "table_name": "product_tracksat",
    "table_stereotype": "sat",
    "satellite_parent_table": "product_hub",
    "compare_criteria": "none"
  }
]
```



```

    },
    {
      "table_name": "product_hub",
      "table_stereotype": "hub",
      "hub_key_column_name": "HK_PRODUCT"
    }
  ]

```

Effectivity satellite on a link {DV-5.3.4}

An effectivity satellite is a satellite on a link without any field mappings. The compiler will detect this and provide this observation in the `is_effectivity_sat` table attribute to the load process.

```

"fields": [
  {
    "field_name": "ORDER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "order_hub"}]
  },
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  }
],
"tables": [
  {
    "table_name": "order_product_sale_esat",
    "table_stereotype": "sat",
    "satellite_parent_table": "order_product_link",
    "driving_keys": ["HK_ORDER"]
  },
  {
    "table_name": "order_product_link",
    "table_stereotype": "lnk",
    "link_key_column_name": "LK_order_product",
    "link_parent_tables": ["order_hub", "product_hub"]
  },
  {
    "table_name": "order_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_ORDER"
  },
  {
    "table_name": "product_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_PRODUCT"
  }
]

```

Effectivity satellite on a link with membership columns {DV-5.3.4}

At least the content of membership start column is created from the source. In this example the membership end column is treated the same way as the load enddate, but refers to the `membership_start` column.

```

"fields": [
  {
    "field_name": "CUSTOMER_ID",
    "field_type": "integer",
    "targets": [{"table_name": "customer_hub"}]
  },
  {
    "field_name": "LOYALTYPROGRAM_NAME",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "loyaltyprogram_hub"}]
  },
  {
    "field_name": "START_DATE",
    "field_type": "date",
    "targets": [{"table_name": "loyaltyprogram_customer_mbsat"}]},
],
"tables": [
  {
    "table_name": "loyaltyprogram_customer_mbsat",
    "table_stereotype": "sat",
    "satellite_parent_table": "loyaltyprogram_customer_link",
    "driving_keys": ["HK_CUSTOMER"],
    "membership_end_columns": [
      {"column_name": "END_DATE",
        "membership_start_column": "START_DATE"}]
  },
  {
    "table_name": "loyaltyprogram_customer_link",
    "table_stereotype": "lnk",
    "link_key_column_name": "LK_LOYALTYPROGRAM_CUSTOMER",
    "link_parent_tables": ["order_hub", "product_hub"],
  },
  {
    "table_name": "customer_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_CUSTOMER"
  },
  {
    "table_name": "loyaltyprogram_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_LOYALTYPROGRAM"
  }
]

```

Normalized record tracking satellites {DV-5.3.5}

Record tracking information must be inserted every time we get it. This is achieved by setting the compare criteria to "none".

```

"fields": [
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "APPERANCE",
    "field_type": "integer",
    "targets": [{"table_name": "product_rectracksat"}]
  }
]

```

```

    },
    "tables": [
      {
        "table_name": "product_rectracksat",
        "table_stereotype": "sat",
        "satellite_parent_table": "product_hub",
        "compare_criteria": "none"
      },
      {
        "table_name": "product_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_PRODUCT"
      }
    ]
  }

```

Normalized record tracking satellites keeping only last record

Record tracking information must be inserted every time we get it, but previous recordings are removed the the history_depth_criteria=0 declaration to keep the table small.

```

"fields": [
  {
    "field_name": "PRODUCT_ID",
    "field_type": "Varchar(20)",
    "targets": [{"table_name": "product_hub"}]
  },
  {
    "field_name": "APPERANCE",
    "field_type": "integer",
    "targets": [{"table_name": "product_rectracksat"}]
  }
],
"tables": [
  {
    "table_name": "product_rectracksat",
    "table_stereotype": "sat",
    "satellite_parent_table": "product_hub",

    "compare_criteria": "none", "history_depth_criteria": "versions", "history_depth_limit": 0
  },
  {
    "table_name": "product_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_PRODUCT"
  }
]

```

"ref" tables

"ref" simple {DV-6.3.1}

By setting is_enddated=false, a load process can only maintain a consistent image of the source data by deleting rows from the ref table, that are missing now.

```

"fields": [
  {
    "field_name": "ISO_3166_ALPHA_2",
    "field_type": "Varchar(2)",
    "targets": [{"table_name": "country_code_ref"}]
  },
  {
    "field_name": "ISO_3166_ALPHA_3",
    "field_type": "Varchar(2)",
    "targets": [{"table_name": "country_code_ref"}]
  },
  {
    "field_name": "ISO_3166_NUMERIC",
    "field_type": "Varchar(3)",
    "targets": [{"table_name": "country_code_ref"}]
  }
],
"tables": [
  {
    "table_name": "country_code_ref",
    "table_stereotype": "ref",
    "is_enddated": "false"
  },
]

```

"ref" historized (hub/sat) {DV-6.3.2}

```

"fields": [
  {
    "field_name": "DAY_DATE",
    "field_type": "DATE",
    "targets": [{"table_name": "calender_hub"}]
  },
  {
    "field_name": "FISCAL_YEAR",
    "field_type": "integer",
    "targets": [{"table_name": "calender_sat"}]
  },
  {
    "field_name": "FISCAL_QUATER",
    "field_type": "integer",
    "targets": [{"table_name": "calender_sat"}]
  }
],
"tables": [
  {
    "table_name": "calender_sat",
    "table_stereotype": "sat",
    "satellite_parent_table": "calender_hub",
    "uses_diff_hash": false
  },
  {
    "table_name": "calender_hub",
    "table_stereotype": "hub",
    "hub_key_column_name": "HK_CALENDER"
  }
]

```

"ref" historized (single table)

This is a non historized reference table extended by an enddate. The enddate is set by the loadprocess, when the specific value combination is not present in the source.

```
"fields": [
  {
    "field_name": "ISO_3166_ALPHA_2",
    "field_type": "Varchar(2)",
    "targets": [{"table_name": "country_code_ref"}]
  },
  {
    "field_name": "ISO_3166_ALPHA_3",
    "field_type": "Varchar(3)",
    "targets": [{"table_name": "country_code_ref"}]
  },
  {
    "field_name": "ISO_3166_NUMERIC",
    "field_type": "Varchar(3)",
    "targets": [{"table_name": "country_code_ref"}]
  }
],
"tables": [
  {
    "table_name": "country_code_ref",
    "table_stereotype": "ref",
    "diff_hash_column_name": "DIFF_COUNTRY_CODE_REF"
  },
]
```

"ref" - code and descriptions {DV-6.3.3}

To achieve the "code and description" pattern, a group column is added as a constant field and the column for the code and description are mapped to the standard naming of the pattern.

A partitioned deletion detection must be declared to modify the default behavior of reference table loading, that would delete all rows, currently not in stage.

```
"fields": [
  {
    "field_name": "GROUP",
    "field_type": "Varchar(20)",
    "fiels_value": "StdOpCode",
    "targets": [{"table_name": "codes_ref"}]
  },
  {
    "field_name": "STANDARD_OPERATIONS_CODE",
    "field_type": "Varchar(20)",
    "targets": [{
      "table_name": "codes_ref",
      "column_name": "CODE"
    }]
  },
  {
    "field_name": "OPERATION_NAME",
    "field_type": "Varchar(255)",
    "targets": [{
      "table_name": "codes_ref",
      "column_name": "DESCRIPTION"
    }]
  }
]
```

```

        ]]
    }
},
"tables": [
    {
        "table_name": "codes_ref",
        "table_stereotype": "ref",
        "is_enddated": "false"
    },
]
"deletion_detection": {
    "procedure": "stage_comparison",
    "partitioning_fields": ["GROUP"],
    "deletion_rules": [
        {"tables_to_cleanup": ["codes_ref"]}
    ]
}

```

"ref" - code and descriptions, historized (hub/sat) {DV-6.3.3.1}

```

"fields": [
    {
        "field_name": "GROUP",
        "field_type": "Varchar(20)",
        "fiels_value": "StdOpCode",
        "targets": [{"table_name": "code_hub"}]
    },
    {
        "field_name": "STANDARD_OPERATIONS_CODE",
        "field_type": "Varchar(20)",
        "targets": [{
            "table_name": "code_hub",
            "column_name": "CODE"
        }]
    },
    {
        "field_name": "OPERATION_NAME",
        "field_type": "Varchar(255)",
        "targets": [{
            "table_name": "code_sat",
            "column_name": "DESCRIPTION"
        }]
    }
],
"tables": [
    {
        "table_name": "code_hub",
        "table_stereotype": "hub",
        "hub_key_column_name": "HK_CODE"
    },
    {
        "table_name": "code_sat",
        "table_stereotype": "sat",
        "satellite_parent_table": "code_hub",
        "uses_diff_hash": false
    }
]
"deletion_detection": {
    "procedure": "stage_comparison",

```

```

        "partitioning_fields": ["GROUP"],
        "deletion_rules": [
            {
                "tables_to_cleanup": ["code_sat"],
                "join_path": ["code_hub"]
            }
        ]
    }
}

```

"ref" - code and descriptions, historized (single table)

This is a non historized reference table extendet by an enddate. The enddate is set by a load run, where the specific value combination of the group is not present in the source. Restrict enddating to the group is done with the "deletion_detection" declaration.

```

"fields": [
    {
        "field_name": "GROUP",
        "field_type": "Varchar(20)",
        "fiels_value": "StdOpCode",
        "targets": [{"table_name": "codes_ref"}]
    },
    {
        "field_name": "STANDARD_OPERATIONS_CODE",
        "field_type": "Varchar(20)",
        "targets": [{
            "table_name": "codes_ref",
            "column_name": "CODE"
        }]
    },
    {
        "field_name": "OPERATION_NAME",
        "field_type": "Varchar(255)",
        "targets": [{
            "table_name": "codes_ref",
            "column_name": "DESCRIPTION"
        }]
    }
],
"tables": [
    {
        "table_name": "codes_ref",
        "table_stereotype": "ref",
        "diff_hash_column_name": "DIFF_CODES_REF"
    }
],
"deletion_detection": {
    "procedure": "stage_comparison",
    "partitioning_fields": ["GROUP"],
    "deletion_rules": [
        {
            "tables_to_cleanup": ["codes_ref"]
        }
    ]
}

```

Snapshot tables

The syntax for snapshot tables (point in time tables, bridge tables) will be added in future versions. From the current perspective this will be a new stereotype "snapshot" with properties to declare the participating links and satellites, the definition of the source for the timeline, the management of the time window and granularity and the structural assets (naming of the columns, columnset for every connected asset (loaddate, key, flag))

Point in time tables {DV-6.1}

to be defined later

Bridge tables {DV-6.2}

to be defined later